

In [9]:

Removing all variables...

```
In [9]: import sys
...: import pandas as pd
...:
...: # a. Getting the data
...: file_path = "C:/Users/ujjwa/OneDrive - Centennial College/Documents/Semester_3/Artificial_Intelligence/Assignments/Logistic_Regression/"
...: file_name = "titanic.csv"
...:
...: df = pd.read_csv(file_path + file_name)
...:
...: df.head()
```

In [10]:

```
...: print df.shape
...: print df.dtypes
...: print df.head(3)
```

	PassengerId	Survived	Pclass	...	Fare	Cabin	Embarked
0	1	0	3	...	7.2500	NaN	S
1	2	1	1	...	71.2833	C85	C
2	3	1	3	...	7.9250	NaN	S

[3 rows x 12 columns]

(891, 12)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 891 entries, 0 to 890

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	PassengerId	891 non-null	int64
1	Survived	891 non-null	int64
2	Pclass	891 non-null	int64
3	Name	891 non-null	object
4	Sex	891 non-null	object
5	Age	714 non-null	float64
6	SibSp	891 non-null	int64
7	Parch	891 non-null	int64
8	Ticket	891 non-null	object
9	Fare	891 non-null	float64
10	Cabin	204 non-null	object
11	Embarked	889 non-null	object

dtypes: float64(2), int64(5), object(5)

memory usage: 83.7+ KB

None

```
In [11]: print df['Sex'].value_counts()
...: print df['Pclass'].value_counts()
['male' 'female']
```

```

In [12]: import pandas as pd

In [13]: df = pd.DataFrame({'Survived': 1, 'Pclass': 1, 'Sex': 'male', 'Age': 22.0, 'SibSp': 1, 'Parch': 0, 'Fare': 7.25, 'Embarked': 'S'})
Out[13]: Text(0, 0.5, 'Passenger Survival (1 = survived)')

In [14]: df = pd.DataFrame({'Survived': 1, 'Pclass': 1, 'Sex': 'male', 'Age': 22.0, 'SibSp': 1, 'Parch': 0, 'Fare': 7.25, 'Embarked': 'S'})
Out[14]: Text(0, 0.5, 'Passenger Survival (1 = survived)')

In [15]: df = pd.DataFrame({'Survived': 1, 'Pclass': 1, 'Sex': 'male', 'Age': 22.0, 'SibSp': 1, 'Parch': 0, 'Fare': 7.25, 'Embarked': 'S'})
Out[15]: Text(0, 0.5, 'Passenger Survival (1 = survived)')

In [16]: df = pd.DataFrame({'Survived': 1, 'Pclass': 1, 'Sex': 'male', 'Age': 22.0, 'SibSp': 1, 'Parch': 0, 'Fare': 7.25, 'Embarked': 'S'})
Out[16]: Text(0, 0.5, 'Passenger Survival (1 = survived)')

In [17]: df = pd.DataFrame({'Survived': 1, 'Pclass': 1, 'Sex': 'male', 'Age': 22.0, 'SibSp': 1, 'Parch': 0, 'Fare': 7.25, 'Embarked': 'S'})
Out[17]: Text(0, 0.5, 'Passenger Survival (1 = survived)')

```

```

value_Sex
value_Embarked
Survived  Pclass  Age  ... Embarked_C Embarked_Q Embarked_S
0         0      3  22.0 ...          0          0          1
1         1      1  38.0 ...          1          0          0
2         1      3  26.0 ...          0          0          1
3         1      1  35.0 ...          0          0          1
4         0      3  35.0 ...          0          0          1
..         ...    ...  ...  ...          ...          ...          ...
886        0      2  27.0 ...          0          0          1
887        1      1  19.0 ...          0          0          1
888        0      3   NaN ...          0          0          1
889        1      1  26.0 ...          1          0          0
890        0      3  32.0 ...          0          1          0

```

[891 rows x 11 columns]

```

In [18]:          = 'Age' .
...:          'Age' = 'Age' .

```

```

In [19]:          = . float

```

```

In [20]:          .
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Survived    891 non-null    float64
1   Pclass      891 non-null    float64
2   Age         891 non-null    float64
3   SibSp       891 non-null    float64
4   Parch       891 non-null    float64
5   Fare        891 non-null    float64
6   Sex_female  891 non-null    float64
7   Sex_male    891 non-null    float64
8   Embarked_C  891 non-null    float64
9   Embarked_Q  891 non-null    float64
10  Embarked_S  891 non-null    float64
dtypes: float64(11)
memory usage: 76.7 KB

```

```

In [21]: def normalize_dataframe
...:     = - . / . - .
...:     return

```

```

In [22]:          =

```

```

In [23]: print          2
Survived  Pclass  Age  ... Embarked_C Embarked_Q Embarked_S
0         0.0    1.0  0.271174 ...          0.0          0.0          1.0
1         1.0    0.0  0.472229 ...          1.0          0.0          0.0

```

[2 rows x 11 columns]

```

In [24]:
...:

In [25]: from
import

In [26]: =

In [27]: = = 'Survived'

In [28]: = 'Survived'

In [29]: = 84

In [30]: = = =0.3 =

In [31]: from import
...: from import
...: from import
...: import as

In [32]: =

In [33]:
Out[33]: LogisticRegression()

In [34]: = . zip .

In [35]: print
0 1
0 Pclass [-2.0024965886472677]
1 Age [-1.671489410236257]
2 SibSp [-1.3543887932123984]
3 Parch [-0.8021837011776475]
4 Fare [0.4348431336707537]
5 Sex_female [1.4779698424175023]
6 Sex_male [-1.4775097824861738]
7 Embarked_C [0.19737518913601654]
8 Embarked_Q [-0.006712720850837989]
9 Embarked_S [-0.3465841109942421]

In [36]: = None
...: = 0

In [37]: = . 0.10 0.50 0.05

In [38]: for in
...: # Splitting the data into training and test sets based on the current
test_size
...: =
=
...:
...: # Creating and fitting the logistic regression model
...: =
...: .

```

```

....:
....:     # 10-fold cross-validation and collect accuracy scores
....:     =                                     =10
....:     = 'accuracy'
....:
....:     # Calculating the minimum, mean, and maximum accuracy scores
....:     = .
....:     = .
....:     = .
....:
....:     # Printing the results for the current test size
....:     print f"Test Size: {           :.0%}"
....:     print f"Minimum Accuracy: {           :.4f}"
....:     print f"Mean Accuracy: {           :.4f}"
....:     print f"Maximum Accuracy: {           :.4f}"
....:     print
....:
....:     # Checking if the current test size has a better mean accuracy
....:     if >
....:         =
....:         =

```

Test Size: 10%
Minimum Accuracy: 0.7125
Mean Accuracy: 0.7991
Maximum Accuracy: 0.9000

Test Size: 15%
Minimum Accuracy: 0.7237
Mean Accuracy: 0.8085
Maximum Accuracy: 0.8947

Test Size: 20%
Minimum Accuracy: 0.7361
Mean Accuracy: 0.8119
Maximum Accuracy: 0.8732

Test Size: 25%
Minimum Accuracy: 0.6866
Mean Accuracy: 0.8159
Maximum Accuracy: 0.8955

Test Size: 30%
Minimum Accuracy: 0.6935
Mean Accuracy: 0.8138
Maximum Accuracy: 0.9194

Test Size: 35%
Minimum Accuracy: 0.7241
Mean Accuracy: 0.8168
Maximum Accuracy: 0.8966

Test Size: 40%
Minimum Accuracy: 0.7037
Mean Accuracy: 0.8204
Maximum Accuracy: 0.8868

Test Size: 45%
 Minimum Accuracy: 0.6531
 Mean Accuracy: 0.8122
 Maximum Accuracy: 0.8980

```
In [39]: print f"Recommended Best Test Size: {          :.0%}"
Recommended Best Test Size: 40%
```

```
In [40]:
          =
          =0.3          =
```

```
In [41]:          =          .          1
```

```
In [42]:          =          > 0.5
```

```
In [43]: from          import
```

```
In [44]:          =
...: print
0.7388059701492538
```

```
In [45]:          =
...: print
[[131 31]
 [ 39 67]]
```

```
In [46]:          =
...: print
```

	precision	recall	f1-score	support
0.0	0.77	0.81	0.79	162
1.0	0.68	0.63	0.66	106
accuracy			0.74	268
macro avg	0.73	0.72	0.72	268
weighted avg	0.74	0.74	0.74	268

```
In [47]:          =          > 0.75
```

```
In [48]:          =
...: print
0.7574626865671642
```

```
In [49]: from          import
```

```
In [50]:
...:          =
...: print
[[158 4]
```

```
[ 61 45]]

In [51]:
....: print f"Accuracy on Training Data: {          :.4f}"
Accuracy on Training Data: 0.8218

In [52]:
....:
....: =
....: print f"Accuracy on Test Data: {          :.4f}"
Accuracy on Test Data: 0.7388

In [53]:
....: if
....:     print "Accuracy on Test Data is higher than Training Data."
....: else
....:     print "Accuracy on Training Data is higher than Test Data."
Accuracy on Training Data is higher than Test Data.

In [54]:
....: =
....: print f"Precision at Threshold 0.5: {          :.4f}"
....:
....: =
....: print f"Recall at Threshold 0.5: {          :.4f}"
....:
....: # Calculating precision score and recall score at threshold 0.75
....: =
....: print f"Precision at Threshold 0.5: {          :.4f}"
....:
....: =
....: print f"Recall at Threshold 0.5: {          :.4f}"
Precision at Threshold 0.5: 0.6837
Recall at Threshold 0.5: 0.6321
Precision at Threshold 0.5: 0.9184
Recall at Threshold 0.5: 0.4245

In [55]:
....: if
....:     print "Accuracy at Threshold 0.5 is higher than at Threshold 0.75."
....: elif
....:     print "Accuracy at Threshold 0.75 is higher than at Threshold 0.5."
....: else
....:     print "Accuracy is the same at both thresholds."
Accuracy at Threshold 0.75 is higher than at Threshold 0.5.

In [56]:
....:
....: if
....:     print "Precision at Threshold 0.5 is higher than at Threshold 0.75."
....: elif
....:     print "Precision at Threshold 0.75 is higher than at Threshold 0.5."
....: else
....:     print "Precision is the same at both thresholds."
```

```
....:
....: # Comparing the recall
....: if                                     >
....:     print "Recall at Threshold 0.5 is higher than at Threshold 0.75."
....: elif                                   <
....:     print "Recall at Threshold 0.75 is higher than at Threshold 0.5."
....: else
....:     print "Recall is the same at both thresholds."
```

Precision at Threshold 0.75 is higher than at Threshold 0.5.

Recall is the same at both thresholds.

In [57]: